

Remove Outliers

May 24, 2019

0.0.1 Remove Ourliers

In this section, We will remove outliers using different approaches. First, we will start by usig RANSAC algorithm which is able to define the inliers and outliers.

```
In [2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.linear_model import RANSACRegressor    # Robust method for regression and
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_val_score
from sklearn.pipeline import Pipeline
from sklearn.model_selection import GridSearchCV
from sklearn.svm import SVR
from keras.models import Sequential
from keras.layers import Dense
from keras.wrappers.scikit_learn import KerasRegressor
from keras.callbacks import EarlyStopping, ModelCheckpoint
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

%matplotlib inline
```

```
In [3]: data = pd.read_csv('reduced_var_data.csv', index_col = 0)
y = data['SalePrice']
x = data.drop(labels = 'SalePrice', axis=1)
print(data.shape)
data.head()
```

(1459, 32)

```
Out[3]:
```

	LotFrontage	LotArea	OverallQual	OverallCond	YearBuilt	YearRemodAdd	\
Id							
1	65.0	8450.0	7.0	5.0	2003.0	2003.0	
2	80.0	9600.0	6.0	8.0	1976.0	1976.0	
3	68.0	11250.0	7.0	5.0	2001.0	2002.0	
4	60.0	9550.0	7.0	5.0	1915.0	1970.0	

5	84.0	14260.0	8.0	5.0	2000.0	2000.0
---	------	---------	-----	-----	--------	--------

	ExterQual	BsmtQual	BsmtCond	BsmtExposure	...	GarageFinish	\
Id					...		
1	3.0	4.0	3.0	1.0	...	2.0	
2	2.0	4.0	3.0	4.0	...	2.0	
3	3.0	4.0	3.0	2.0	...	2.0	
4	2.0	3.0	4.0	1.0	...	1.0	
5	3.0	4.0	3.0	3.0	...	2.0	

	GarageCars	GarageArea	OpenPorchSF	MoSold	MSZoning_RL	\
Id						
1	2.0	548.0	61.0	2.0	1.0	
2	2.0	460.0	0.0	5.0	1.0	
3	2.0	608.0	42.0	9.0	1.0	
4	3.0	642.0	35.0	2.0	1.0	
5	3.0	836.0	84.0	12.0	1.0	

	Neighborhood_Crawfor	MasVnrArea_209.0	MasVnrArea_428.0	SalePrice
Id				
1	0.0	0.0	0.0	208500
2	0.0	0.0	0.0	181500
3	0.0	0.0	0.0	223500
4	1.0	0.0	0.0	140000
5	0.0	0.0	0.0	250000

[5 rows x 32 columns]

0.1 Find Outliers using RANSACRegressor

In [4]: *# Define outliers in the data set*

```
ransac = RANSACRegressor(LinearRegression(),
                           max_trials=1000,
                           min_samples=1000,
                           loss='absolute_loss',
                           random_state=42)
```

```
ransac.fit(x, y)
```

Out [4]: RANSACRegressor(base_estimator=LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False), is_data_valid=None, is_model_valid=None, loss='absolute_loss', max_skips=inf, max_trials=1000, min_samples=1000, random_state=42, residual_threshold=None, stop_n_inliers=inf, stop_probability=0.99, stop_score=inf)


In [5]: inlier_mask = ransac.inlier_mask_
outlier_mask = np.logical_not(inlier_mask)

```
In [6]: outlier_mask.sum(), inlier_mask.sum()
```

```
Out[6]: (152, 1307)
```

Now we removed the outliers from the dataset. The original data contains 1459 examples and the data without any outliers have 1307 examples. In other words, we removed 152 examples.

```
In [7]: plt.hist(y[inlier_mask], bins = 30);
```



0.2 SVR on data without outliers

```
In [8]: x_train, x_test, y_train, y_test = train_test_split(x[inlier_mask], y[inlier_mask], test_size=0.2, random_state=42)
```

```
In [9]: pipe = Pipeline(steps= [('ss', StandardScaler()), ('clf', SVR(gamma='scale'))])
```

```
param_grid = {
    'clf__C': [0.1, 0.5, 1.0, 1.5, 10, 100, 150, 1000],
    'clf__kernel': ['linear', 'rbf', 'sigmoid', 'poly']
}
```

```
search = GridSearchCV(pipe, param_grid, cv=5, iid=False, scoring='neg_mean_absolute_error', return_train_score=False)
```

```
search.fit(x, y) # Here I am using the whole training data
print("Best parameter (CV score=%0.3f):" % search.best_score_)
print(search.best_params_)
```

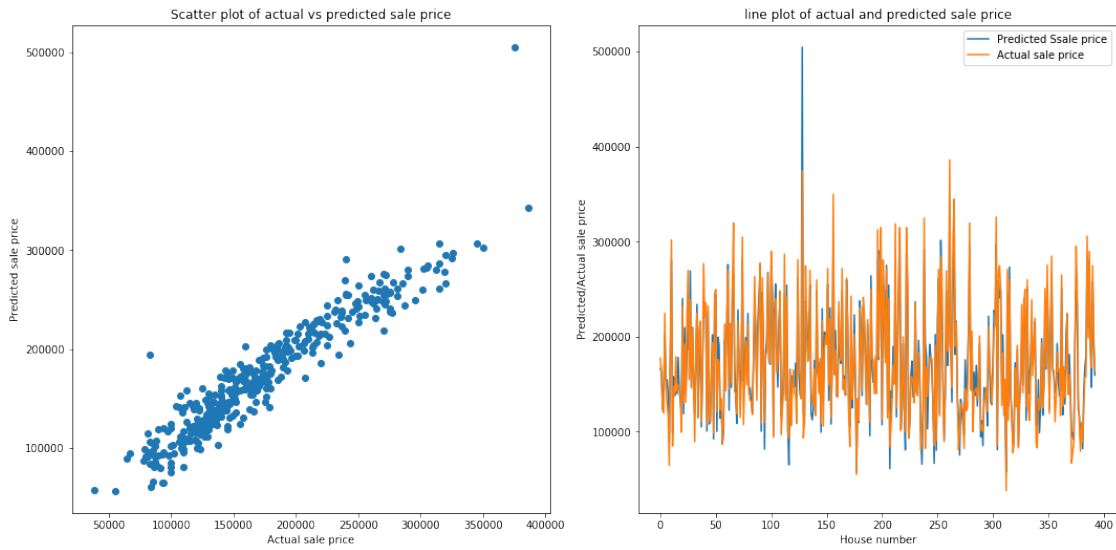
```
Best parameter (CV score=-19636.219):  
{'clf__C': 1000, 'clf__kernel': 'linear'}
```

```
In [10]: ss = StandardScaler()  
         ss.fit(x_train)  
         x_train = ss.transform(x_train)  
         x_test = ss.transform(x_test)  
  
         best_svr = SVR(kernel='linear', gamma='scale', C = 1000)  
         best_svr.fit(x_train, y_train)  
         y_pred = best_svr.predict(x_test)  
         print(mean_absolute_error(y_test, y_pred))  
         print('R2_score = {}'.format(r2_score(y_test, y_pred)))  
  
14062.252899828172  
R2_score = 0.9027249959650371
```

```
In [11]: print(best_svr.score(x_test, y_test))  
  
0.9027249959650371
```

```
In [12]: y_pred = best_svr.predict(x_test)  
         fig = plt.figure(figsize=(15,8))  
         fig.suptitle('SVR Results without Outliers')  
         plt.subplot(121)  
         plt.scatter(y_test.values, y_pred)  
         plt.xlabel('Actual sale price')  
         plt.ylabel('Predicted sale price')  
         plt.title('Scatter plot of actual vs predicted sale price')  
         plt.subplot(122)  
         plt.plot((y_pred), label='Predicted Ssale price')  
         plt.plot((y_test.values), label='Actual sale price')  
         plt.xlabel('House number')  
         plt.ylabel('Predicted/Actual sale price')  
         plt.title('line plot of actual and predicted sale price')  
         plt.legend()  
         plt.tight_layout()  
         fig.subplots_adjust(top=0.88)
```


SVR Results without Outliers



```
In [13]: print("Corrolation between true and predicted value using SVR on the actual sale price is",
              format(np.corrcoef(y_test,y_pred)[0][1]))
```

Corrolation between true and predicted value using SVR on the actual sale price is 0.950250546

```
In [14]: resid = y_test - y_pred
         mean_resid = resid.mean()
         std_resid = resid.std()
         plt.plot(y_test,y_test-y_pred,'.')
         plt.xlabel('y')
         plt.ylabel('y - y_pred');
         plt.title('std resid = {:.3f}'.format(std_resid));
```



0.3 Neural Networks

In [15]: `x_train, x_test, y_train, y_test = train_test_split(x[inlier_mask], y[inlier_mask], t`

```
ss = StandardScaler()
ss.fit(x_train)
x_train = ss.transform(x_train)
x_test = ss.transform(x_test)
```

In [16]: `seed = 42`

```
np.random.seed(seed)

n_feat = x.shape[1]
mdl = Sequential()
mdl.add(Dense(units=256, input_dim = n_feat, activation = 'relu'))
mdl.add(Dense(units=1, activation='linear'))

mdl.compile(loss='mean_squared_error', optimizer='adam', metrics= ['mse', 'mae'])
monitor = EarlyStopping(monitor= 'val_loss', min_delta=1e-3,
                        patience = 10, verbose=1, mode = 'auto')
history = mdl.fit(x_train, y_train, validation_data = (x_test, y_test),
                  callbacks=[monitor], batch_size= 64, epochs=5000, verbose=0)

y_pred = mdl.predict(x_test)
```

```

MSEscore = (mean_squared_error(y_pred, y_test))
print('Score MSE = {}'.format(MSEscore))

MAEscore = (mean_absolute_error(y_pred, y_test))
print('Score MAE = {}'.format(MAEscore))
print('R2_score = {}'.format(r2_score(y_test, y_pred)))
mdl.summary()

```

```

Epoch 04124: early stopping
Score MSE = 380631885.78712326
Score MAE = 13701.82104802799
R2_score = 0.90407417909799

```

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 256)	8192
dense_2 (Dense)	(None, 1)	257

```

Total params: 8,449
Trainable params: 8,449
Non-trainable params: 0

```

```


In [17]: plt.figure()
         plt.plot(history.epoch, history.history['mean_squared_error'], label = 'Train Error')
         plt.plot(history.epoch, history.history['val_mean_squared_error'], label = 'Val Error')
         plt.legend()

```


```

Out[17]: <matplotlib.legend.Legend at 0x2256aa6ab00>

```



```
In [18]: y_pred = mdl.predict(x_test)
fig = plt.figure(figsize=(15,8))
fig.suptitle('Neural Networks Results without Outliers')
plt.subplot(121)
plt.scatter((y_pred), (y_test))
plt.xlabel('Actual sale price')
plt.ylabel('Predicted sale price')
plt.title('Scatter plot of actual vs predicted sale price')
plt.subplot(122)
plt.subplots_adjust(bottom=0.25)
plt.plot((y_pred), label='Predicted Ssale price')
plt.plot((y_test.values), label='Actual sale price')
plt.xlabel('House number')
plt.ylabel('Predicted/Actual sale price')
plt.title('line plot of actual and predicted sale price')
plt.legend()
plt.tight_layout()
fig.subplots_adjust(top=0.88)
```

```
In [19]: print("Corrolation between true and predicted value using NN on the actual sale price
              format(np.corrcoef(y_test,y_pred.squeeze())[0][1]))
```

Corrolation between true and predicted value using NN on the actual sale price is 0.9516045835

```
In [20]: resid = y_test - y_pred.squeeze()
          mean_resid = resid.mean()
          std_resid = resid.std()
          plt.plot(y_test,resid,'.')
          plt.xlabel('y')
          plt.ylabel('y - y_pred');
          plt.title('std resid = {:.3f}'.format(std_resid));
```

