

remove outliers and log chang y to improve performance

May 24, 2019

I try to remove some outliers in X for better performance. Also use log transfer y. Fit in Ridge and Lasso with clean data. Use grid search to improve performance.

```
In [1]: import numpy as np
import pandas as pd
from matplotlib import pyplot as plt
import seaborn as sns
```

```
%matplotlib inline
```

```
In [2]: clean_data = pd.read_csv('clean_data.csv')
clean_data.head()
```

```
Out[2]:
```

	Id	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	YearBuilt	\
0	1	60	65.0	8450	7	5	2003	
1	2	20	80.0	9600	6	8	1976	
2	3	60	68.0	11250	7	5	2001	
3	4	70	60.0	9550	7	5	1915	
4	5	60	84.0	14260	8	5	2000	

	YearRemodAdd	ExterQual	ExterCond	...	SaleType_ConLI	\
0	2003	3	2	...		0
1	1976	2	2	...		0
2	2002	3	2	...		0
3	1970	2	2	...		0
4	2000	3	2	...		0

	SaleType_ConLw	SaleType_New	SaleType_Oth	SaleType_WD	\
0	0	0	0	1	
1	0	0	0	1	
2	0	0	0	1	
3	0	0	0	1	
4	0	0	0	1	

	SaleCondition_AdjLand	SaleCondition_Alloca	SaleCondition_Family	\
0	0	0	0	
1	0	0	0	
2	0	0	0	

3	0	0	0
4	0	0	0

	SaleCondition_Normal	SaleCondition_Partial
0	1	0
1	1	0
2	1	0
3	0	0
4	1	0

[5 rows x 534 columns]

```
In [3]: clean_data.columns
```

```
Out[3]: Index(['Id', 'MSSubClass', 'LotFrontage', 'LotArea', 'OverallQual',
              'OverallCond', 'YearBuilt', 'YearRemodAdd', 'ExterQual', 'ExterCond',
              ...,
              'SaleType_ConLI', 'SaleType_ConLw', 'SaleType_New', 'SaleType_Oth',
              'SaleType_WD', 'SaleCondition_AdjLand', 'SaleCondition_Alloca',
              'SaleCondition_Family', 'SaleCondition_Normal',
              'SaleCondition_Partial'],
              dtype='object', length=534)
```

```
In [4]: y=clean_data['SalePrice']
        clean_data=clean_data.drop(['SalePrice'],axis=1) #drop the y in clean_data
        X = clean_data.iloc[:,1:] # drop the id column
        X.shape
```


```
Out[4]: (1459, 532)
```

```
In [5]: y.describe()
```

```
Out[5]: count      1459.000000
        mean      180930.394791
        std       79468.964025
        min       34900.000000
        25%      129950.000000
        50%      163000.000000
        75%      214000.000000
        max       755000.000000
        Name: SalePrice, dtype: float64
```

```
In [6]: plt.hist(y,bins=15)
        plt.title('Distribution of y')
        # Basically the first and last 9 bins have little samples,
        #it's hard to predict these small groups of data,so drop it first.
```

```
Out[6]: Text(0.5,1,'Distribution of y')
```



```
In [7]: from sklearn.model_selection import train_test_split
        from sklearn.linear_model import Lasso
        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=
        # scaling
        from sklearn.preprocessing import StandardScaler
        scaler = StandardScaler()
        X_train_sc = scaler.fit_transform(X_train)
        X_test_sc = scaler.transform(X_test)
```

```
/Users/fanwenyu/anaconda3/lib/python3.6/site-packages/sklearn/preprocessing/data.py:625: DataC
    return self.partial_fit(X, y)
/Users/fanwenyu/anaconda3/lib/python3.6/site-packages/sklearn/base.py:462: DataConversionWarni
    return self.fit(X, **fit_params).transform(X)
/Users/fanwenyu/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:8: DataConversionW
```

```
In [8]: clf = Lasso(random_state=42)
        clf.fit(X_train_sc, y_train)
        y_test_pred = clf.predict(X_test_sc)
        y_train_pred = clf.predict(X_train_sc)
        from sklearn.metrics import mean_squared_error
        from math import sqrt
        mse = mean_squared_error(y_test, y_test_pred)
        sqrt(mse)
```



```

ConvergenceWarning)
/Users/fanwenyu/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:
ConvergenceWarning)
/Users/fanwenyu/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:
ConvergenceWarning)
/Users/fanwenyu/anaconda3/lib/python3.6/site-packages/sklearn/model_selection/_search.py:841:
DeprecationWarning)

```

```

In [10]: mse_test = mean_squared_error(y_test, y_test_pred)
         sqrt(mse_test)

```

```

Out[10]: 26802.87448312158

```

```


In [11]: plt.scatter(y_test, y_test_pred)
         plt.title('Performance of lasso with sheared data(RMSE=33665.39)')
         plt.xlabel('y_test')
         plt.ylabel('y_test_pred')

```

```

Out[11]: Text(0,0.5,'y_test_pred')

```



```

In [12]: clf.best_score_
         # R2

```

```

Out[12]: 0.8037315767569586

```

0.1 Remove outliers in X

```
In [13]: y_true_comb = np.hstack([y_train,y_test])  
        len(y_true_comb)
```

```
Out[13]: 1459
```

```
In [14]: X_true_comb = np.vstack([X_train_sc,X_test_sc])  
        X_true_comb = pd.DataFrame(X_true_comb,columns=X.columns)
```


```
In [15]: y_pred = clf.predict(X_true_comb)
```

```
In [16]: error =y_true_comb-y_pred  
        error
```

```
Out[16]: array([-14028.48554686,  18869.90785598,  31935.99323894, ...,  
               20869.21771686, -35026.38340765, -29035.94844561])
```

```
In [17]: plt.hist(y_true_comb-y_pred,bins = 30)  
        plt.title('Distribution of errors')
```

```
Out[17]: Text(0.5,1,'Distribution of errors')
```



```
In [18]: strange_X = X_true_comb[error<-60000]  
        strange_X
```

```

Out[18]:
MSSubClass  LotFrontage  LotArea  OverallQual  OverallCond  YearBuilt  \
700      0.052093   -1.673021  0.148034   -0.090550    1.328141   0.176179
1064    -0.881004    1.196449  0.191906    1.376930   -0.494482   1.197285
1094    -0.881004    0.259479 -0.175261   -0.824290   -0.494482  -0.647293
1152    -0.881004    0.903646  0.075925    1.376930   -0.494482   1.098468
1167    -0.881004    1.723495  0.265211    1.376930   -0.494482   1.131407
1173    -0.881004   -1.673021  2.028113   -0.090550   -2.317105  -0.482598
1218     0.285367   -0.209006 -0.194760   -1.558029   -1.405794  -1.701337
1299    -0.881004    0.230199  1.331579   -0.824290    2.239452  -0.120270

YearRemodAdd  ExterQual  ExterCond  BsmtQual  ...  \
700      0.492087  -0.692465  2.661743  0.577925  ...
1064     1.165509  2.756348 -0.217113  1.689151  ...
1094    -1.576280 -2.416871 -0.217113 -0.533302  ...
1152     1.021204  1.031941 -0.217113  1.689151  ...
1167     1.021204  1.031941 -0.217113  0.577925  ...
1173    -0.469944  1.031941 -0.217113 -0.533302  ...
1218    -1.672484 -0.692465 -3.095969 -0.533302  ...
1299     0.876900 -0.692465 -0.217113 -0.533302  ...

SaleType_ConLI  SaleType_ConLw  SaleType_New  SaleType_Oth  SaleType_WD  \
700      -0.031311   -0.044302   -0.309020   -0.031311    0.393672
1064     -0.031311   -0.044302    3.236033   -0.031311   -2.540188
1094     -0.031311   -0.044302   -0.309020   -0.031311    0.393672
1152     -0.031311   -0.044302    3.236033   -0.031311   -2.540188
1167     -0.031311   -0.044302   -0.309020   -0.031311    0.393672
1173     -0.031311   -0.044302   -0.309020   -0.031311    0.393672
1218     -0.031311   -0.044302   -0.309020   -0.031311    0.393672
1299     -0.031311   -0.044302   -0.309020   -0.031311    0.393672

SaleCondition_AdjLand  SaleCondition_Alloca  SaleCondition_Family  \
700      -0.062715   -0.076885   -0.126176
1064     -0.062715   -0.076885   -0.126176
1094     -0.062715   13.006409   -0.126176
1152     -0.062715   -0.076885   -0.126176
1167     -0.062715   -0.076885   -0.126176
1173     -0.062715   13.006409   -0.126176
1218     -0.062715   -0.076885   -0.126176
1299     -0.062715   -0.076885   -0.126176

SaleCondition_Normal  SaleCondition_Partial
700      0.470417   -0.310918
1064     -2.125775    3.216278
1094     -2.125775   -0.310918
1152     -2.125775    3.216278
1167      0.470417   -0.310918
1173     -2.125775   -0.310918
1218      0.470417   -0.310918

```

1299 -2.125775 3.216278

[8 rows x 532 columns]

In [19]: strange_X = X_true_comb[error>80000]
strange_X

```
Out [19]:
```

	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	YearBuilt	\
390	-0.881004	1.401411	0.442724	2.84441	-0.494482	1.197285	
446	0.052093	1.459972	0.301082	2.11067	-0.494482	1.197285	
491	0.402004	0.962207	1.134291	2.84441	3.150764	-2.623626	
554	-0.881004	0.083797	-0.232562	1.37693	0.416829	1.164346	
1316	0.052093	1.372131	0.266775	2.11067	-0.494482	1.197285	
1379	0.052093	3.011828	0.460384	2.84441	-0.494482	0.802018	

	YearRemodAdd	ExterQual	ExterCond	BsmtQual	...	\
390	1.117407	2.756348	-0.217113	1.689151	...	
446	1.165509	2.756348	-0.217113	1.689151	...	
491	0.395884	1.031941	2.661743	-0.533302	...	
554	1.069306	1.031941	-0.217113	0.577925	...	
1316	1.165509	2.756348	-0.217113	1.689151	...	
1379	0.540189	1.031941	-0.217113	1.689151	...	

	SaleType_ConLI	SaleType_ConLw	SaleType_New	SaleType_Oth	SaleType_WD	\
390	-0.031311	-0.044302	-0.309020	-0.031311	0.393672	
446	-0.031311	-0.044302	3.236033	-0.031311	-2.540188	
491	-0.031311	-0.044302	-0.309020	-0.031311	0.393672	
554	-0.031311	-0.044302	3.236033	-0.031311	-2.540188	
1316	-0.031311	-0.044302	3.236033	-0.031311	-2.540188	
1379	-0.031311	-0.044302	-0.309020	-0.031311	0.393672	

	SaleCondition_AdjLand	SaleCondition_Alloca	SaleCondition_Family	\
390	-0.062715	-0.076885	-0.126176	
446	-0.062715	-0.076885	-0.126176	
491	-0.062715	-0.076885	-0.126176	
554	-0.062715	-0.076885	-0.126176	
1316	-0.062715	-0.076885	-0.126176	
1379	-0.062715	-0.076885	-0.126176	

	SaleCondition_Normal	SaleCondition_Partial
390	0.470417	-0.310918
446	-2.125775	3.216278
491	0.470417	-0.310918
554	-2.125775	3.216278
1316	-2.125775	3.216278
1379	-2.125775	-0.310918


[6 rows x 532 columns]


```
In [20]: # identify outliers of X according to its feature
```

```
cm = plt.cm.get_cmap('PiYG')
sc = plt.scatter(X_true_comb.iloc[:,1],y_pred , c=error, cmap=cm) # LotFrontage:
                                                                    #Linear feet of s

plt.colorbar(sc)
plt.title('LotFrontage is large in outliers')
plt.xlabel('LotFrontage')
plt.ylabel('y_test_pred')
```

```
Out[20]: Text(0,0.5,'y_test_pred')
```




```
In [21]: # remove according to LotFrontage.
```

```
X_filded = X_true_comb[X_true_comb['LotFrontage'] < 3]
y_filded = y_true_comb[X_true_comb['LotFrontage'] < 3]
error_filded = error[X_true_comb['LotFrontage'] < 3]
cm = plt.cm.get_cmap('PiYG')
sc = plt.scatter(X_filded.iloc[:,1],y_filded , c=error_filded, cmap=cm) # LotFrontage
                                                                    #Linear feet of s

plt.colorbar(sc)
plt.title('LotFrontage is large in outliers')
plt.xlabel('LotFrontage')
plt.ylabel('y_test_pred')
```

```
Out[21]: Text(0,0.5,'y_test_pred')
```



```
In [22]: X_train, X_test, y_train, y_test = train_test_split(X_filted, y_filted, test_size=0.3)
         clf = Lasso(random_state=42)
         clf.fit(X_train, y_train)
         y_test_pred = clf.predict(X_test)
         y_train_pred = clf.predict(X_train)
         from sklearn.metrics import mean_squared_error
         mse = mean_squared_error(y_test, y_test_pred)
         sqrt(mse)
```

```
/Users/fanwenyu/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:
ConvergenceWarning)
```

Out[22]: 35317.15079537594

```
In [23]: parameters = { 'max_iter':[1000,2000,5000], 'alpha':[1, 10,100,1000,10000]}
         ls = Lasso(random_state=42)
         clf = GridSearchCV(ls, parameters, cv=5)
         clf.fit(X_train, y_train)
         y_test_pred = clf.predict(X_test)
         y_train_pred = clf.predict(X_train)
         mse = mean_squared_error(y_test, y_test_pred)
         sqrt(mse)
```

```

/Users/fanwenyu/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:
  ConvergenceWarning)
/Users/fanwenyu/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:
  ConvergenceWarning)
/Users/fanwenyu/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:
  ConvergenceWarning)
/Users/fanwenyu/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:
  ConvergenceWarning)
/Users/fanwenyu/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:
  ConvergenceWarning)
/Users/fanwenyu/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:
  ConvergenceWarning)
/Users/fanwenyu/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:
  ConvergenceWarning)
/Users/fanwenyu/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:
  ConvergenceWarning)
/Users/fanwenyu/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:
  ConvergenceWarning)
/Users/fanwenyu/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:
  ConvergenceWarning)
/Users/fanwenyu/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:
  ConvergenceWarning)
/Users/fanwenyu/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:
  ConvergenceWarning)
/Users/fanwenyu/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:
  ConvergenceWarning)

```

Out [23]: 30570.689624112565

In [24]: `y.shape[0]-y_filted.shape[0]` *# 7 data points removed*

Out [24]: 7


```

In [25]: plt.scatter(y_test, y_test_pred)
         plt.title('Performance of lasso with sheared data according to LotFrontage(RMSE=30570)')
         plt.xlabel('y_test')
         plt.ylabel('y_test_pred')

```

Out [25]: Text(0,0.5,'y_test_pred')

Performance of lasso with sheared data according to LotFrontage(RMSE=30570.68)



```
In [26]: # combine data for some EDA and find more outliers
```

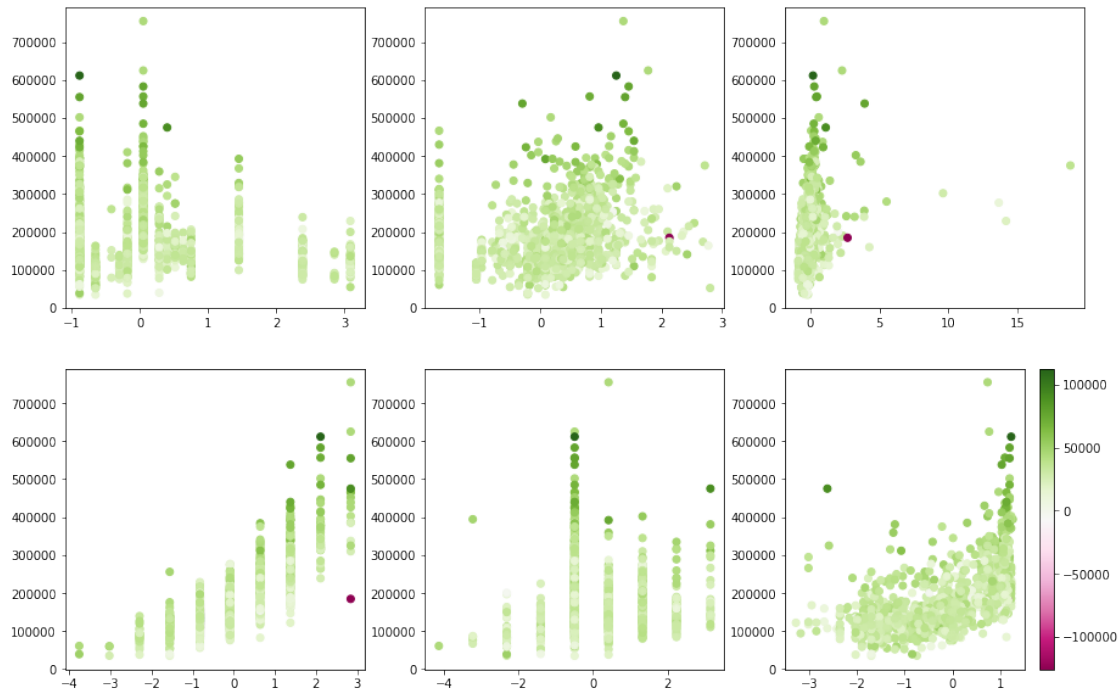
```
y_filded = np.hstack([y_train,y_test])
X_true_comb = np.vstack([X_train,X_test])
X_filded = pd.DataFrame(X_true_comb,columns=X.columns)
y_pred = clf.predict(X_filded)
error_filded =y_filded-y_pred
error_filded
```

```
Out[26]: array([ 5.21606840e+03, -9.47065676e+03,  5.37606815e+03, ...,
                -1.64250998e+02, -1.62202384e+01,  1.63888821e+04])
```

```
In [27]: i = 0
```

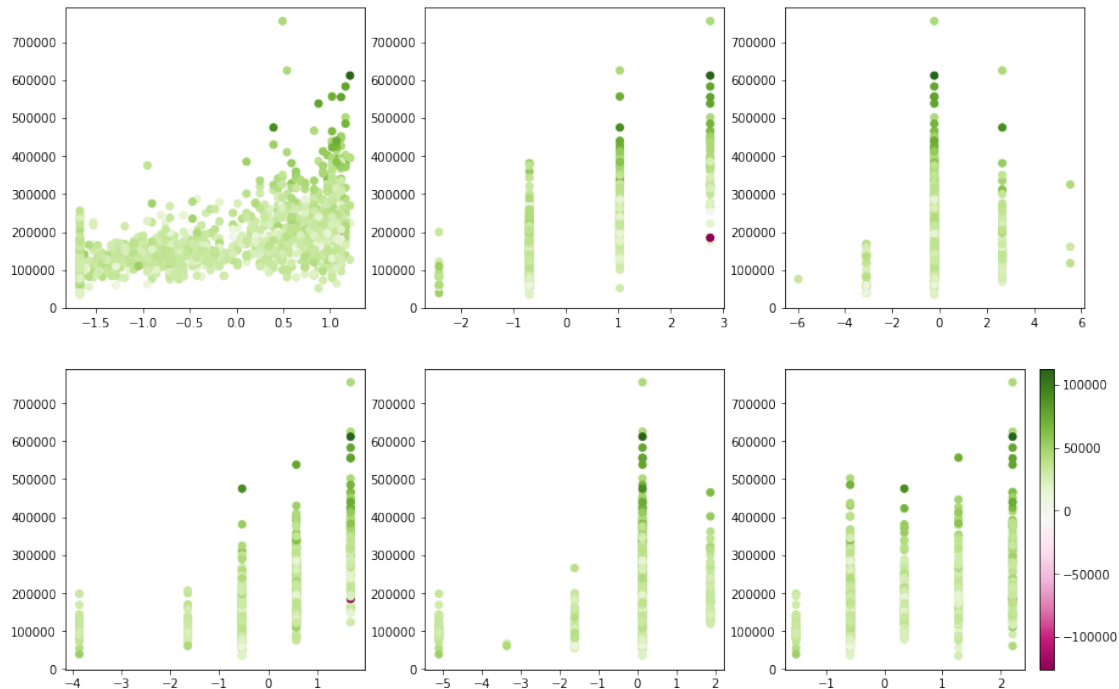
```
fig, axes = plt.subplots(2, 3, figsize=(15,10))
axes[0, 0].scatter(X_filded.iloc[:,i+0],y_filded , c=error_filded, cmap=cm)
axes[0, 1].scatter(X_filded.iloc[:,i+1],y_filded , c=error_filded, cmap=cm)
axes[0, 2].scatter(X_filded.iloc[:,i+2],y_filded , c=error_filded, cmap=cm)
axes[1, 0].scatter(X_filded.iloc[:,i+3],y_filded , c=error_filded, cmap=cm)
axes[1, 1].scatter(X_filded.iloc[:,i+4],y_filded , c=error_filded, cmap=cm)
axes[1, 2].scatter(X_filded.iloc[:,i+5],y_filded , c=error_filded, cmap=cm)
plt.colorbar(sc)
```

```
Out[27]: <matplotlib.colorbar.Colorbar at 0x1a257a4b00>
```



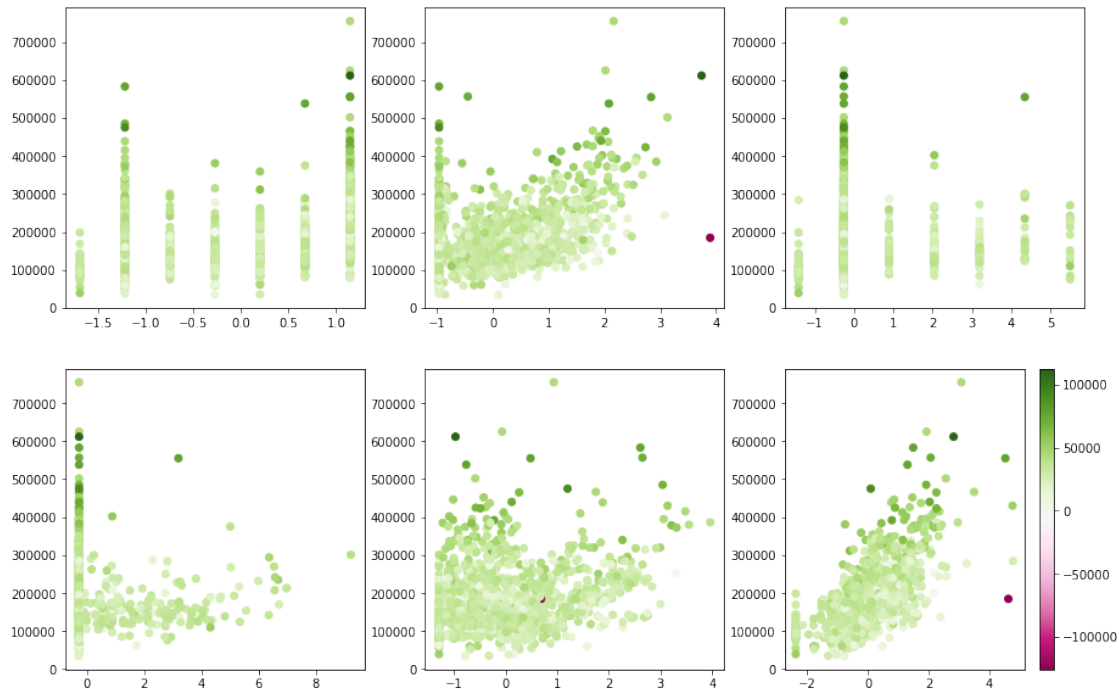
```
In [28]: i = 6
fig, axes = plt.subplots(2, 3, figsize=(15,10))
axes[0, 0].scatter(X_filded.iloc[:,i+0],y_filded , c=error_filded, cmap=cm)
axes[0, 1].scatter(X_filded.iloc[:,i+1],y_filded , c=error_filded, cmap=cm)
axes[0, 2].scatter(X_filded.iloc[:,i+2],y_filded , c=error_filded, cmap=cm)
axes[1, 0].scatter(X_filded.iloc[:,i+3],y_filded , c=error_filded, cmap=cm)
axes[1, 1].scatter(X_filded.iloc[:,i+4],y_filded , c=error_filded, cmap=cm)
axes[1, 2].scatter(X_filded.iloc[:,i+5],y_filded , c=error_filded, cmap=cm)
plt.colorbar(sc)
```

```
Out[28]: <matplotlib.colorbar.Colorbar at 0x1a13b234e0>
```



```
In [29]: i = 12
fig, axes = plt.subplots(2, 3, figsize=(15,10))
axes[0, 0].scatter(X_filded.iloc[:,i+0],y_filded , c=error_filded, cmap=cm)
axes[0, 1].scatter(X_filded.iloc[:,i+1],y_filded , c=error_filded, cmap=cm)
axes[0, 2].scatter(X_filded.iloc[:,i+2],y_filded , c=error_filded, cmap=cm)
axes[1, 0].scatter(X_filded.iloc[:,i+3],y_filded , c=error_filded, cmap=cm)
axes[1, 1].scatter(X_filded.iloc[:,i+4],y_filded , c=error_filded, cmap=cm)
axes[1, 2].scatter(X_filded.iloc[:,i+5],y_filded , c=error_filded, cmap=cm)
plt.colorbar(sc)
```

```
Out[29]: <matplotlib.colorbar.Colorbar at 0x1a13df9198>
```



```
In [30]: # find some outliers in feature BsmtFinSF1
import heapq
print(heapq.nlargest(3, X_filted.iloc[:,13]))
```

```
[3.8974066763514945, 3.7424466786490984, 3.131215576600758]
```

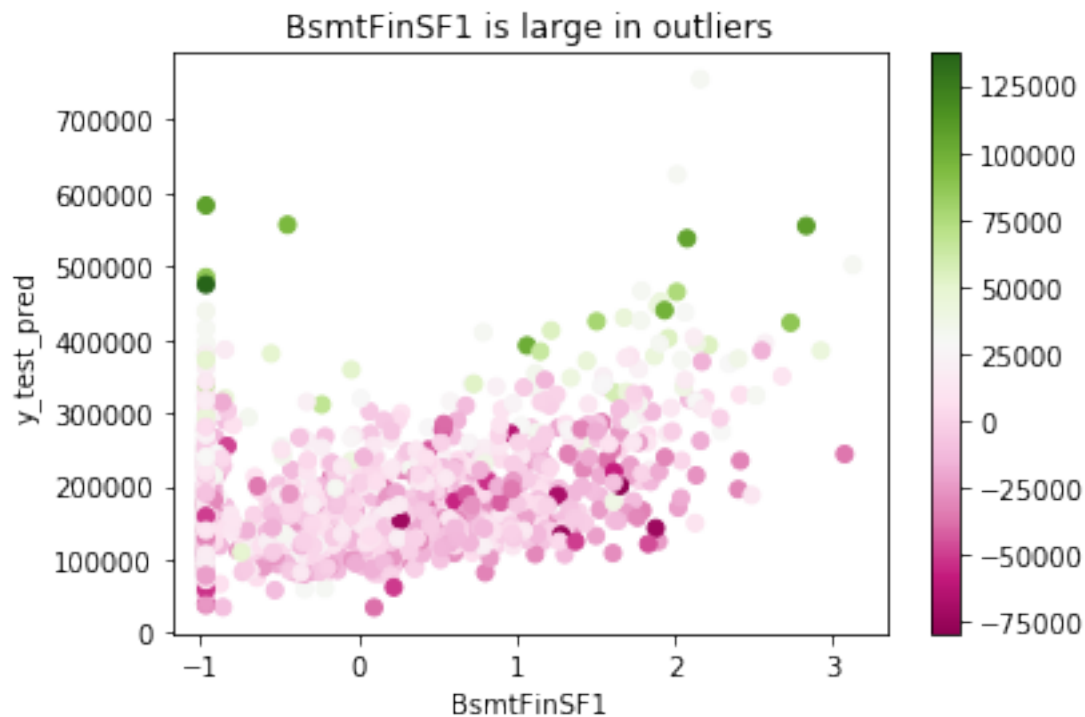
```
In [31]: X_filted.columns[13]
```

```
Out[31]: 'BsmtFinSF1'
```

```
In [33]: # remove according to BsmtFinSF1.
X_filted2 = X_filted[X_filted['BsmtFinSF1'] < 3.6]
y_filted2 = y_filted[X_filted['BsmtFinSF1'] < 3.6]
error_filted2 = error_filted[X_filted['BsmtFinSF1'] < 3.6]
cm = plt.cm.get_cmap('PiYG')
sc = plt.scatter(X_filted2.iloc[:,13],y_filted2 , c=error_filted2, cmap=cm) # LotFrontage
#Linear feet of street frontage

plt.colorbar(sc)
plt.title('BsmtFinSF1 is large in outliers')
plt.xlabel('BsmtFinSF1')
plt.ylabel('y_test_pred')
```

```
Out[33]: Text(0,0.5,'y_test_pred')
```



```
In [45]: X_train, X_test, y_train, y_test = train_test_split(X_filted2, y_filted2, test_size=0.2)

parameters = { 'max_iter':[1000,2000,5000], 'alpha':[1, 10,100,1000,10000]}

ls = Lasso(random_state=42)

clf = GridSearchCV(ls, parameters, cv=5)

clf.fit(X_train,y_train)

y_test_pred = clf.predict(X_test)

y_train_pred = clf.predict(X_train)

mse = mean_squared_error(y_test, y_test_pred)

sqrt(mse)

/Users/fanwenyu/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:140: ConvergenceWarning:
  /Users/fanwenyu/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:140: ConvergenceWarning)
/Users/fanwenyu/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:140: ConvergenceWarning:
  /Users/fanwenyu/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:140: ConvergenceWarning)
/Users/fanwenyu/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:140: ConvergenceWarning:
  /Users/fanwenyu/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:140: ConvergenceWarning)
```



```

/Users/fanwenyu/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:
ConvergenceWarning)
/Users/fanwenyu/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:
ConvergenceWarning)
/Users/fanwenyu/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:
ConvergenceWarning)
/Users/fanwenyu/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:
ConvergenceWarning)
/Users/fanwenyu/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:
ConvergenceWarning)
/Users/fanwenyu/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:
ConvergenceWarning)
/Users/fanwenyu/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:
ConvergenceWarning)
/Users/fanwenyu/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:
ConvergenceWarning)
/Users/fanwenyu/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:
ConvergenceWarning)
/Users/fanwenyu/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:
ConvergenceWarning)

```

Out[45]: 25183.600856565372

0.2 After remove the outliers in X, the root mean squared error dropped to 25183.60


```
In [46]: y_filted.shape[0]-y_filted2.shape[0] # 2 data points removed
```

Out[46]: 2

```
In [47]: plt.scatter(y_test, y_test_pred)
plt.title('Performance of lasso with sheared data according to BsmtFinSF1(RMSE=25183.60)')
plt.xlabel('y_test')
plt.ylabel('y_test_pred')
```

Out[47]: Text(0,0.5,'y_test_pred')

Performance of lasso with sheared data according to BsmtFinSF1(RMSE=25183.60)



0.3 fit into Ridge

```
In [48]: parameters = { 'max_iter':[1000,2000,5000], 'alpha':[1, 10,100,1000,10000]}
          from sklearn.linear_model import Ridge
          ls = Ridge(random_state=42)
          clf = GridSearchCV(ls, parameters, cv=5)
          clf.fit(X_train,y_train)
          y_test_pred = clf.predict(X_test)
          y_train_pred = clf.predict(X_train)
```

```
In [49]: mse_test = mean_squared_error(y_test, y_test_pred)
          sqrt(mse_test)
```


```
Out[49]: 26037.908039527552
```

```
In [50]: clf.best_score_
```

```
Out[50]: 0.8699814937005034
```

```
In [51]: plt.scatter(y_test, y_test_pred)
          plt.title('Performance of Ridge with sheared data(RMSE=26037.91)')
          plt.xlabel('y_test')
          plt.ylabel('y_test_pred')
```

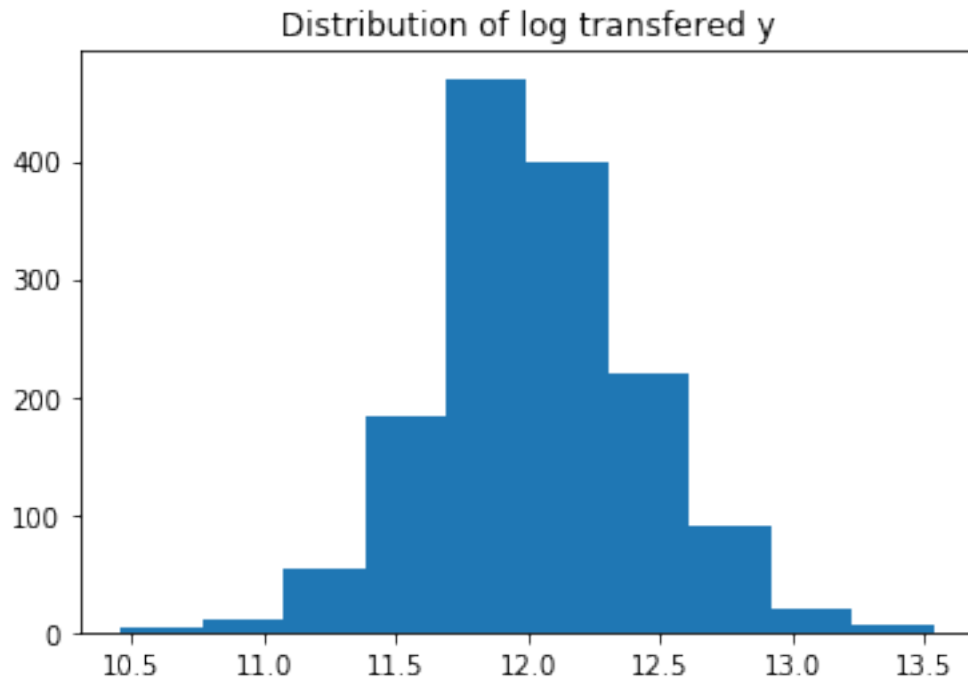
```
Out[51]: Text(0,0.5,'y_test_pred')
```



0.4 log transfered y

```
In [52]: # log transfered y
y_log = np.log(y)
plt.hist(y_log)
plt.title('Distribution of log transfered y')
```

```
Out[52]: Text(0.5,1,'Distribution of log transfered y')
```



0.5 fit into lasso

In [58]: `X_train, X_test, y_train, y_test = train_test_split(X, y_log, test_size=0.3, random_s`


```
clf = Lasso(random_state=42)
clf.fit(X_train,y_train)
y_test_pred = clf.predict(X_test)
y_train_pred = clf.predict(X_train)

mse = mean_squared_error(y_test, y_test_pred)
sqrt(mse)
```

Out[58]: 0.18036476431273618

```
In [63]: parameters = { 'max_iter':[100,500,1000], 'alpha':[0.001,0.01,0.1,1, 10,100,1000]}
ls = Lasso(random_state=42)
clf = GridSearchCV(ls, parameters, cv=5)
clf.fit(X_train,y_train)
y_test_pred = clf.predict(X_test)
y_train_pred = clf.predict(X_train)
```

```
/Users/fanwenyu/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:
ConvergenceWarning)
/Users/fanwenyu/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:
ConvergenceWarning)
```


0.6 fit into Ridge

```
In [69]: parameters = { 'max_iter':[10,50,100,500], 'alpha':[0.001,0.01,0.1,1, 10,100,1000]}
          rg = Ridge(random_state=42)
          clf = GridSearchCV(rg, parameters, cv=5)
          clf.fit(X_train,y_train)
          y_test_pred = clf.predict(X_test)
          y_train_pred = clf.predict(X_train)
```

```
/Users/fanwenyu/anaconda3/lib/python3.6/site-packages/sklearn/model_selection/_search.py:841: DeprecationWarning
```

```
In [70]: plt.scatter(y_test, y_test_pred)
          plt.title('Performance of Ridge with log transfer y(RMSE=0.134803)')
          plt.xlabel('y_test')
          plt.ylabel('y_test_pred')
```

```
Out[70]: Text(0,0.5,'y_test_pred')
```



```
In [71]: mse_test = mean_squared_error(y_test, y_test_pred)
         sqrt(mse_test)
```

```
Out[71]: 0.1375474209000953
```

```
In [72]: clf.best_score_
```

```
Out[72]: 0.6234102765983629
```