

```
[20] import numpy as np
import pandas as pd
from matplotlib import pyplot as plt
import seaborn as sns
from scipy import stats
from math import ceil

%matplotlib inline
```

1. Load data

```
[52] df = pd.read_csv('../train.csv', index_col=0)
df.head()
```

| | MSSubClass | MSZoning | LotFrontage | LotArea | Street | Alley | Id |
|----|------------|----------|-------------|---------|--------|-------|----|
| Id | | | | | | | |
| 1 | 60 | RL | 65.0 | 8450 | Pave | NaN | 1 |
| 2 | 20 | RL | 80.0 | 9600 | Pave | NaN | 1 |
| 3 | 60 | RL | 68.0 | 11250 | Pave | NaN | 1 |
| 4 | 70 | RL | 60.0 | 9550 | Pave | NaN | 1 |
| 5 | 60 | RL | 84.0 | 14260 | Pave | NaN | 1 |

5 rows × 80 columns

```
[53] df.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 1460 entries, 1 to 1460
Data columns (total 80 columns):
MSSubClass      1460 non-null int64
MSZoning        1460 non-null object
LotFrontage     1201 non-null float64
LotArea         1460 non-null int64
Street          1460 non-null object
Alley           91 non-null object
LotShape        1460 non-null object
LandContour     1460 non-null object
Utilities       1460 non-null object
LotConfig       1460 non-null object
LandSlope       1460 non-null object
Neighborhood    1460 non-null object
Condition1      1460 non-null object
```

| | | | |
|--------------|------|----------|---------|
| Condition2 | 1460 | non-null | object |
| BldgType | 1460 | non-null | object |
| HouseStyle | 1460 | non-null | object |
| OverallQual | 1460 | non-null | int64 |
| OverallCond | 1460 | non-null | int64 |
| YearBuilt | 1460 | non-null | int64 |
| YearRemodAdd | 1460 | non-null | int64 |
| RoofStyle | 1460 | non-null | object |
| RoofMatl | 1460 | non-null | object |
| Exterior1st | 1460 | non-null | object |
| Exterior2nd | 1460 | non-null | object |
| MasVnrType | 1452 | non-null | object |
| MasVnrArea | 1452 | non-null | float64 |
| ExterQual | 1460 | non-null | object |
| ExterCond | 1460 | non-null | object |
| Foundation | 1460 | non-null | object |

2. Clean Data

2.1 Columns with NaN Values

```
[54] cols_with_na = df.isnull().sum()
cols_with_na = cols_with_na[cols_with_na>0]
print(cols_with_na.sort_values(ascending=False))
```

| | |
|--------------|------|
| PoolQC | 1453 |
| MiscFeature | 1406 |
| Alley | 1369 |
| Fence | 1179 |
| FireplaceQu | 690 |
| LotFrontage | 259 |
| GarageYrBlt | 81 |
| GarageType | 81 |
| GarageFinish | 81 |
| GarageQual | 81 |
| GarageCond | 81 |
| BsmtFinType2 | 38 |
| BsmtExposure | 38 |
| BsmtFinType1 | 37 |
| BsmtCond | 37 |
| BsmtQual | 37 |
| MasVnrArea | 8 |
| MasVnrType | 8 |
| Electrical | 1 |

dtype: int64

2.2 Meaningful NaN Values

```
[83] df.Alley = df.Alley.fillna(value = 'NoAlley')
df.BsmtCond = df.BsmtCond.fillna(value = 'NoBsmt')
df.BsmtQual = df.BsmtQual.fillna(value = 'NoBsmt')
df.BsmtExposure = df.BsmtExposure.fillna(value= 'NoBsmt')
df.BsmtFinType1 = df.BsmtFinType1.fillna(value= 'NoBsmt')
df.BsmtFinType2 = df.BsmtFinType2.fillna(value= 'NoBsmt')
df.LotFrontage = df.LotFrontage.fillna(value = 0)
df.FireplaceQu = df.FireplaceQu.fillna(value = 'NoFireplace')
df.GarageType = df.GarageType.fillna(value = 'NoGarage')
df.GarageCond = df.GarageCond.fillna(value = 'NoGarage')
df.GarageFinish = df.GarageFinish.fillna(value = 'NoGarage')
df.GarageYrBlt = df.GarageYrBlt.fillna(value = 0)
df.GarageQual = df.GarageQual.fillna(value = 'NoGarage')

df.PoolQC = df.PoolQC.fillna(value = 'NoPool')
df.Fence = df.Fence.fillna(value = 'NoFence')
df.MiscFeature = df.MiscFeature.fillna(value = 'NoMisc')
df.MasVnrType = df.MasVnrType.fillna(value = 'noMas')
df.MasVnrArea = df.MasVnrArea.fillna(value = 'noMas')

df.info()
```


```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1460 entries, 1 to 1460
Data columns (total 80 columns):
MSSubClass      1460 non-null category
MSZoning        1460 non-null category
LotFrontage     1460 non-null float64
LotArea         1460 non-null int64
Street          1460 non-null category
Alley           1460 non-null category
LotShape        1460 non-null category
LandContour     1460 non-null category
Utilities       1460 non-null category
LotConfig       1460 non-null category
LandSlope       1460 non-null category
Neighborhood    1460 non-null category
Condition1      1460 non-null category
Condition2      1460 non-null category
BldgType        1460 non-null category
HouseStyle      1460 non-null category
OverallQual     1460 non-null int64
OverallCond     1460 non-null int64
YearBuilt       1460 non-null int64
YearRemodAdd    1460 non-null int64
RoofStyle       1460 non-null category
RoofMatl        1460 non-null category
Exterior1st     1460 non-null category
Exterior2nd     1460 non-null category
MasVnrType      1460 non-null category
MasVnrArea      1460 non-null category
ExterQual       1460 non-null int64
ExterCond       1460 non-null int64
Foundation      1460 non-null category
```

2.3 Distribution of SalePrice

```
[56] plt.figure(figsize=(12,4))
plt.subplot(1,2,1)
_ = sns.distplot(df.SalePrice.dropna() , fit=stats.norm);
plt.subplot(1,2,2)
_ = stats.probplot(df.SalePrice.dropna(), plot=plt)
```

/Users/changyaohua/anaconda3/lib/python3.7/site-packages/scipy/stats/stats.py:1713: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will be interpreted as an array index, `arr[np.array(seq)]`, which will result either in an error or a different result.


```
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```



2.4 Log Transform SalePrice

```
[57] #Log Transform SalePrice to improve normality
sp = df.SalePrice
df.SalePrice = np.log(sp)

plt.figure(figsize=(12,4))
plt.subplot(1,2,1)
_ = sns.distplot(sp.dropna() , fit=stats.norm);
plt.subplot(1,2,2)
_ = stats.probplot(sp.dropna(), plot=plt)
```



3. Exploratory Data Analysis (EDA)

Different types of features will need to be treated differently when digging deeper in to the data. Here I identify three types of features:

- Numeric-discrete: Numeric features with less than 13 unique values, such as month of the year, or the numeric scales created above.
- Numeric-continuous: Numeric features with 13 or more unique values, such as areas, the year a property was built etc.
- Categorical: The remaining non-numeric features.

```
[58] # extract names of numeric columns
dtypes = df.dtypes
cols_numeric = dtypes[dtypes != object].index.tolist()

# MSubClass should be treated as categorical
cols_numeric.remove('MSubClass')

# choose any numeric column with less than 13 values to be
# "discrete". 13 chosen to include months of the year.
# other columns "continuous"
col_nunique = dict()

for col in cols_numeric:
    col_nunique[col] = df[col].nunique()

col_nunique = pd.Series(col_nunique)

cols_discrete = col_nunique[col_nunique < 13].index.tolist()
cols_continuous = col_nunique[col_nunique >= 13].index.tolist()

print(len(cols_numeric), 'numeric columns, of which',
      len(cols_continuous), 'are continuous and',
      len(cols_discrete), 'are discrete.')
```

35 numeric columns, of which 21 are continuous and 14 are discrete.

```
[59] # extract names of categorical columns
cols_categ = dtypes[~dtypes.index.isin(cols_numeric)].index.tolist()

for col in cols_categ:
    df[col] = df[col].astype('category')

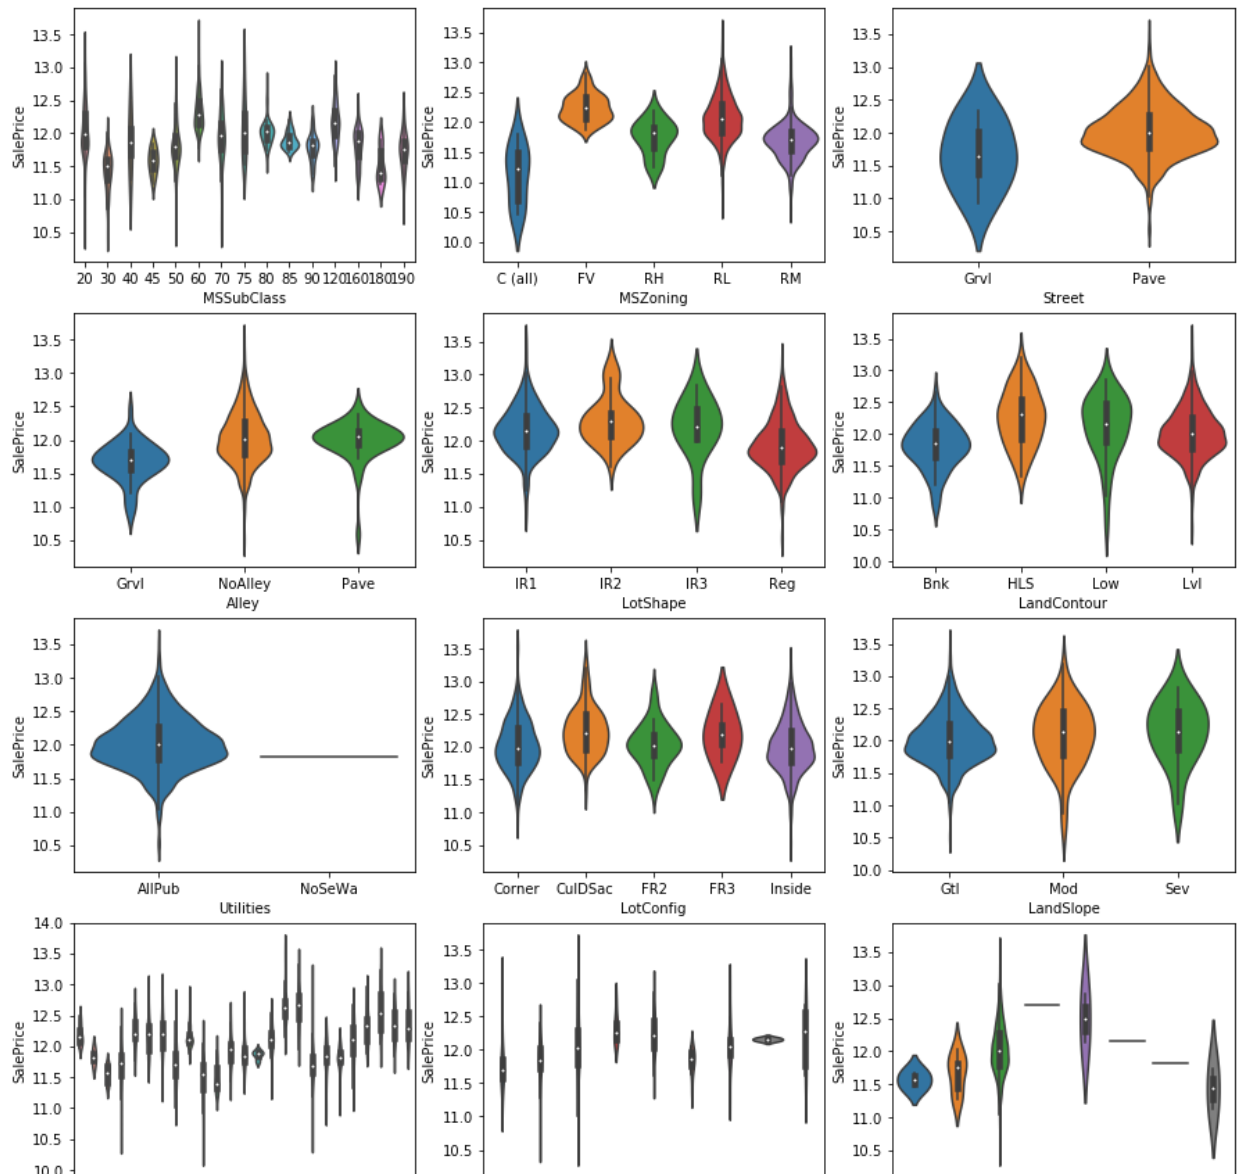
print(len(cols_categ), 'categorical columns.')
```

45 categorical columns.

3.1 Distribution of SalePrice in Categorical Variables


```
[60] # plot categorical variables
fcols = 3
frows = ceil(len(cols_categ)/fcols)
plt.figure(figsize=(15, 4*frows))


for i, col in enumerate(cols_categ):
    plt.subplot(frows, fcols, i+1)
    _ = sns.violinplot(df[col], df['SalePrice'])
```



```
[61] #Neighbourhood
plt.figure(figsize=(25,5))
sns.violinplot(x='Neighborhood',y='SalePrice',data=df)
plt.xticks(rotation=45);

#Exterior1st
plt.figure(figsize=(25,5))
sns.violinplot(x='Exterior1st',y='SalePrice',data=df)
plt.xticks(rotation=45);
```





```
[ 62] df.BsmtCond = df.BsmtCond.map({'Ex':5 , 'Gd':4 , 'TA':3 , 'Fa':2 , 'Po':1 ,
df.BsmtQual = df.BsmtQual.map({'Ex':5 , 'Gd':4 , 'TA':3 , 'Fa':2 , 'Po':1 ,
df.BsmtExposure = df.BsmtExposure.map({'Gd':4 , 'Av':3 , 'Mn':2 , 'No':1 , 'N
df.BsmtFinType1 = df.BsmtFinType1.map({'GLQ':6 , 'ALQ':5 , 'BLQ':4 , 'Rec':3 , 'L
df.BsmtFinType2 = df.BsmtFinType2.map({'GLQ':6 , 'ALQ':5 , 'BLQ':4 , 'Rec':3 , 'L

df.GarageType = df.GarageType.map({'2Types':4 , 'Attchd': 5 , 'Basment':3
                                'CarPort' :1, 'Detchd':2 , 'NoGara

df.GarageCond = df.GarageCond.map({'NoGarage':0 , 'Po':1 , 'Fa':2 , 'TA':3 ,
df.GarageQual = df.GarageQual.map({'NoGarage':0 , 'Po':1 , 'Fa':2 , 'TA':3 ,
df.GarageFinish = df.GarageFinish.map({'Fin':3 , 'RFn':2 , 'Unf':1 , 'NoGara
df.PavedDrive = df.PavedDrive.map({'Y':2 , 'P':1 , 'N':0 })
```

```
[ 63] df.ExterCond = df.ExterCond.map({"Ex":4 , 'Gd':3 , 'TA':2 , 'Fa':1 , 'Po':0})
df.ExterQual = df.ExterQual.map({"Ex":4 , 'Gd':3 , 'TA':2 , 'Fa':1 , 'Po':0})

df.CentralAir = df.CentralAir.map({'Y':1 , 'N':0})
df.HeatingQC = df.HeatingQC.map({"Ex":4 , 'Gd':3 , 'TA':2 , 'Fa':1 , 'Po':0})
df.FireplaceQu = df.FireplaceQu.map({"Ex":5 , 'Gd':4 , 'TA':3 , 'Fa':2 , 'Po':1 ,

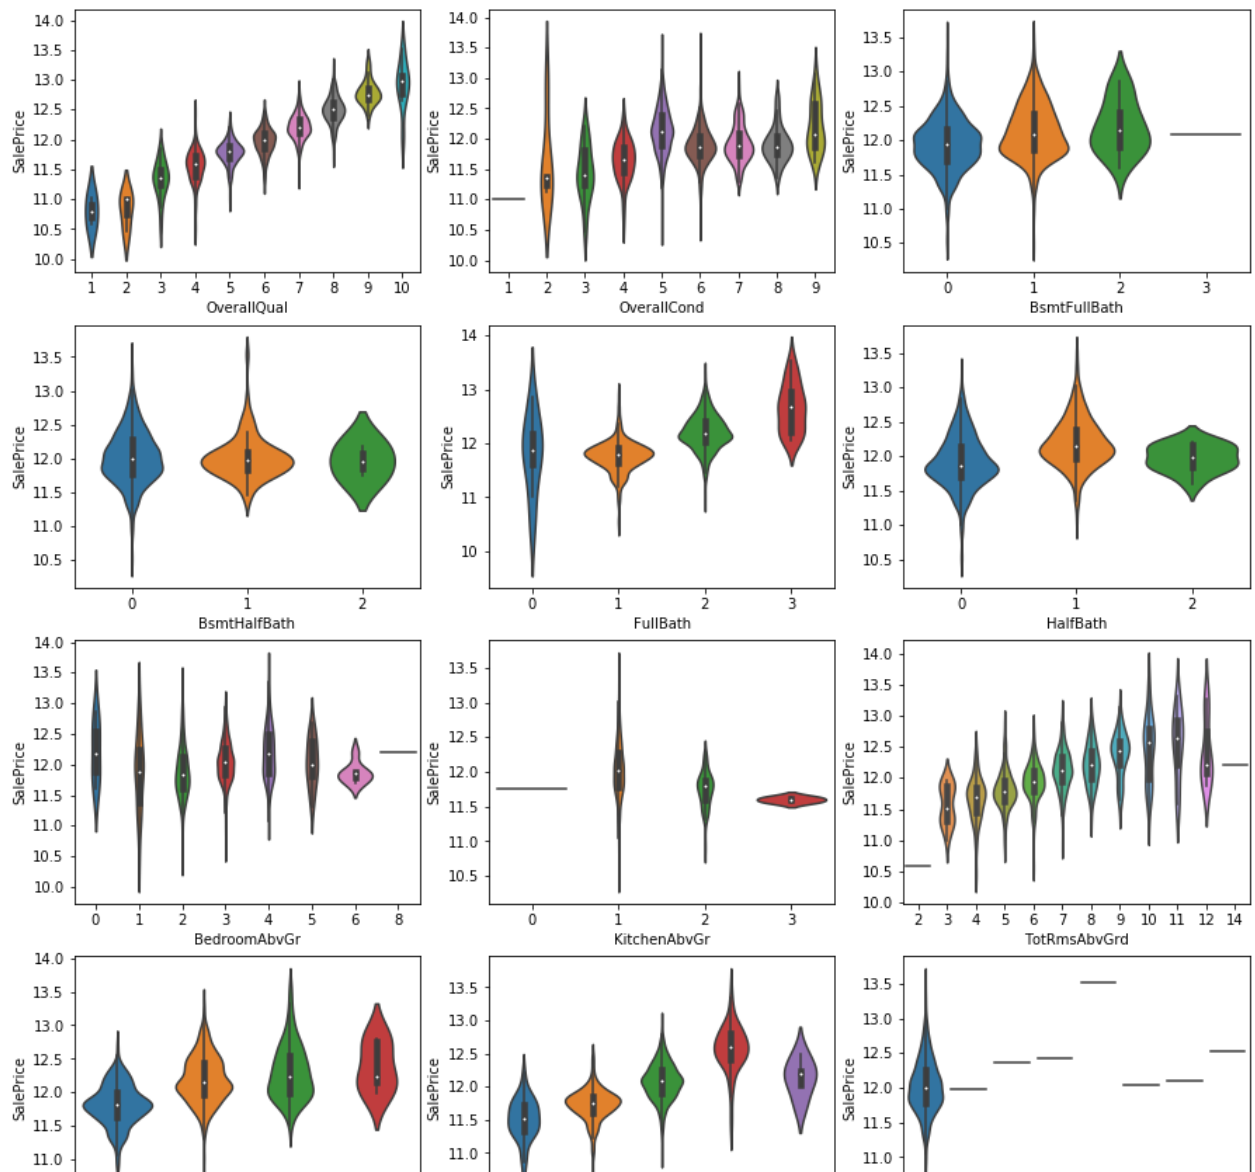
df.KitchenQual = df.KitchenQual.map({"Ex":4 , 'Gd':3 , 'TA':2 , 'Fa':1 , 'Po':0})

df.PoolQC = df.PoolQC.map({"Ex":4 , 'Gd':3 , 'TA':2 , 'Fa':1 , 'NoPool':0})
df.Fence = df.Fence.map({'GdPrv':4 , 'MnPrv':3 , 'GdWo':2 , 'MnWw':1 , 'N
```

3.2 Distribution of SalePrice in Discrete Numeric Features

```
[ 64] fcols = 3
frows = ceil(len(cols_discrete)/fcols)
plt.figure(figsize=(15,4*frows))

for i,col in enumerate(cols_discrete):
    plt.subplot(frows,fcols,i+1)
    sns.violinplot(df[col],df['SalePrice'])
```

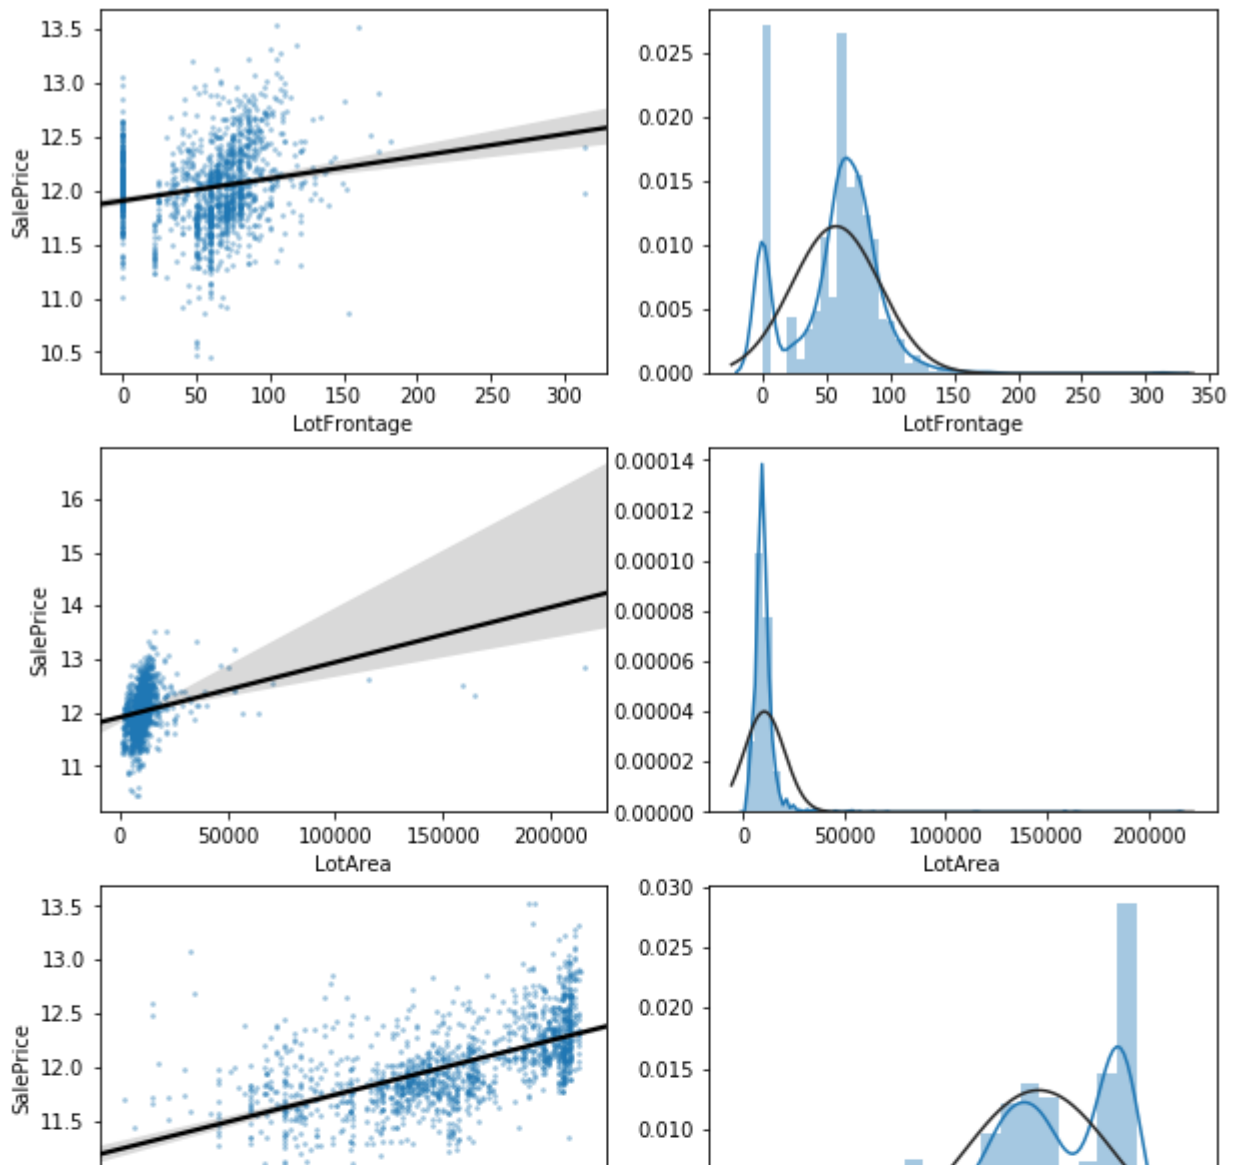
3.3 Distribution of Continuous Variables and Effect on SalePrice

```
[65] fcols = 2
frows = len(cols_continuous)
plt.figure(figsize=(5*fcols,4*frows))

i=0
for col in cols_continuous:
    i+=1
    ax=plt.subplot(frows,fcols,i)
    sns.regplot(x=col, y='SalePrice', data=df, ax=ax,
                scatter_kws={'marker': '.', 's': 3, 'alpha': 0.3},
                line_kws={'color': 'k'});
    plt.xlabel(col)
    plt.ylabel('SalePrice')

    i+=1
    ax=plt.subplot(frows,fcols,i)
    sns.distplot(df[col].dropna() , fit=stats.norm)
```

```
plt.xlabel(col)
```




3.4 Correlation Between Numeric Features

```
[66] # correlation between numeric variables
df_corr = df[cols_numeric].corr(method='spearman').abs()

# order columns and rows by correlation with SalePrice
df_corr = df_corr.sort_values('SalePrice',axis=0,ascending=False).sort_va

ax=plt.figure(figsize=(20,16)).gca()
sns.heatmap(df_corr,ax=ax,square=True)
```

<matplotlib.axes._subplots.AxesSubplot at 0x1a24f63e48>



```
[67] df.info()
```

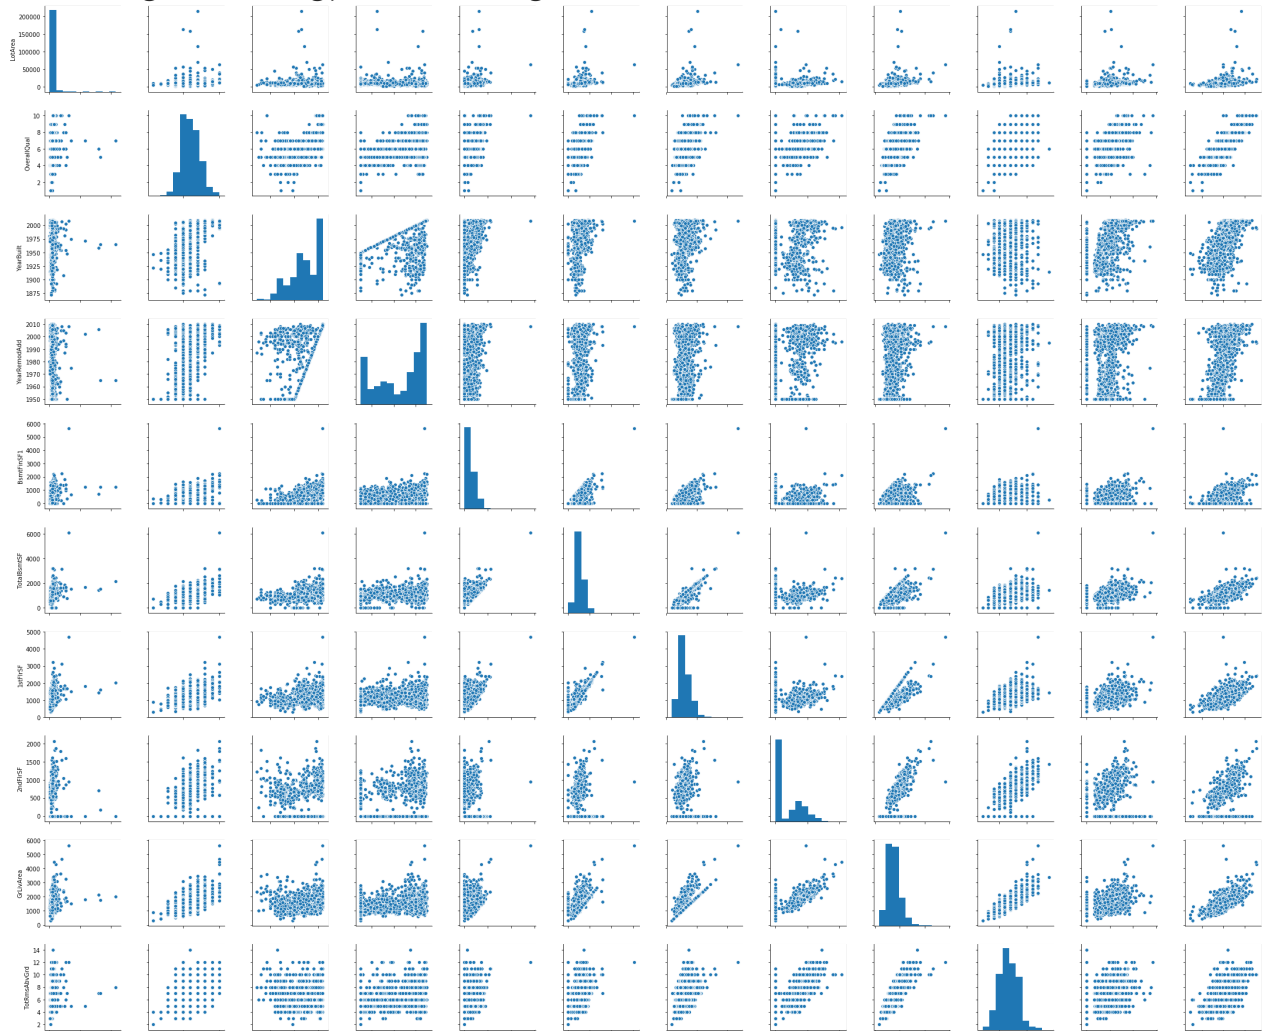
```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1460 entries, 1 to 1460
Data columns (total 80 columns):
MSSubClass      1460 non-null category
MSZoning        1460 non-null category
LotFrontage     1460 non-null float64
LotArea         1460 non-null int64
Street          1460 non-null category
Alley           1460 non-null category
LotShape        1460 non-null category
LandContour     1460 non-null category
Utilities       1460 non-null category
LotConfig       1460 non-null category
LandSlope       1460 non-null category
Neighborhood    1460 non-null category
Condition1      1460 non-null category
Condition2      1460 non-null category
BldgType        1460 non-null category
HouseStyle      1460 non-null category
OverallQual     1460 non-null int64
OverallCond     1460 non-null int64
YearBuilt       1460 non-null int64
YearRemodAdd    1460 non-null int64
RoofStyle       1460 non-null category
RoofMatl        1460 non-null category
Exterior1st     1460 non-null category
Exterior2nd     1460 non-null category
MasVnrType      1460 non-null category
MasVnrArea      1460 non-null category
ExterQual       1460 non-null int64
ExterCond       1460 non-null int64
```

```
[68] # cols = ['LotArea', 'OverallQual', 'YearBuilt', 'YearRemodAdd', 'MasVnrA',
cols = ['LotArea', 'OverallQual', 'YearBuilt', 'YearRemodAdd', 'BsmtFinSF
```


```
[69] sns.pairplot(df[cols], size=2.5)
plt.tight_layout()
plt.show()
```

/Users/changyaohua/anaconda3/lib/python3.7/site-packages/seaborn/axisgrid.py:2065: UserWarning: The `size` parameter has been renamed to `height`; please update your code.

warnings.warn(msg, UserWarning)



```
[70] fig = plt.figure(figsize=(15,8))
cm = np.corrcoef(df[cols].values.T)
sns.set(font_scale=0.8)
hm = sns.heatmap(cm, cbar=True, annot=True, square=True, fmt='.2f', annot
plt.show()
```



```
[71] df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1460 entries, 1 to 1460
Data columns (total 80 columns):
MSSubClass      1460 non-null category
MSZoning        1460 non-null category
LotFrontage     1460 non-null float64
LotArea         1460 non-null int64
Street          1460 non-null category
Alley           1460 non-null category
LotShape        1460 non-null category
LandContour     1460 non-null category
Utilities       1460 non-null category
LotConfig       1460 non-null category
LandSlope       1460 non-null category
Neighborhood    1460 non-null category
Condition1      1460 non-null category
Condition2      1460 non-null category
BldgType        1460 non-null category
HouseStyle      1460 non-null category
OverallQual     1460 non-null int64
OverallCond     1460 non-null int64
YearBuilt       1460 non-null int64
YearRemodAdd    1460 non-null int64
RoofStyle       1460 non-null category
```

| | |
|-------------|------------------------|
| RoofMatl | 1460 non-null category |
| Exterior1st | 1460 non-null category |
| Exterior2nd | 1460 non-null category |
| MasVnrType | 1460 non-null category |
| MasVnrArea | 1460 non-null category |
| ExterQual | 1460 non-null int64 |
| ExterCond | 1460 non-null int64 |
| Foundation | 1460 non-null category |

3.5 Identify and Remove Outliers

```
[78] from sklearn.metrics import make_scorer
from sklearn.linear_model import Ridge
```

```
[75] # metric for evaluation
def rmse(y_true, y_pred):
    diff = y_pred - y_true
    sum_sq = sum(diff**2)
    n = len(y_pred)

    return np.sqrt(sum_sq/n)

# scorer to be used in sklearn model fitting
rmse_scorer = make_scorer(rmse, greater_is_better=False)
```

```
[76] # function to detect outliers based on the predictions of a model
def find_outliers(model, X, y, sigma=3):

    # predict y values using model
    try:
        y_pred = pd.Series(model.predict(X), index=y.index)
    # if predicting fails, try fitting the model first
    except:
        model.fit(X,y)
        y_pred = pd.Series(model.predict(X), index=y.index)

    # calculate residuals between the model prediction and true y values
    resid = y - y_pred
    mean_resid = resid.mean()
    std_resid = resid.std()

    # calculate z statistic, define outliers to be where |z|>sigma
    z = (resid - mean_resid)/std_resid
    outliers = z[abs(z)>sigma].index

    # print and plot the results
    print('R2=',model.score(X,y))
    print('rmse=',rmse(y, y_pred))
    print('-----')
```

```

print('mean of residuals:',mean_resid)
print('std of residuals:',std_resid)
print('-----')

print(len(outliers),'outliers:')
print(outliers.tolist())

plt.figure(figsize=(15,5))
ax_131 = plt.subplot(1,3,1)
plt.plot(y,y_pred, '.')
plt.plot(y.loc[outliers],y_pred.loc[outliers],'ro')
plt.legend(['Accepted','Outlier'])
plt.xlabel('y')
plt.ylabel('y_pred');

ax_132=plt.subplot(1,3,2)
plt.plot(y,y-y_pred, '.')
plt.plot(y.loc[outliers],y.loc[outliers]-y_pred.loc[outliers],'ro')
plt.legend(['Accepted','Outlier'])
plt.xlabel('y')
plt.ylabel('y - y_pred');

ax_133=plt.subplot(1,3,3)
z.plot.hist(bins=50,ax=ax_133)
z.loc[outliers].plot.hist(color='r',bins=50,ax=ax_133)
plt.legend(['Accepted','Outlier'])
plt.xlabel('z')

plt.savefig('outliers.png')

return outliers

```

```

[84] d_df = pd.get_dummies(df, drop_first= True)
d_df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 1460 entries, 1 to 1460
Columns: 546 entries, LotFrontage to SaleCondition_Partial
dtypes: float64(3), int64(50), uint8(493)
memory usage: 1.3 MB

```

```

[85] y = d_df.SalePrice
X = d_df.drop('SalePrice',axis=1)

# find and remove outliers using a Ridge model
outliers = find_outliers(Ridge(), X, y)

# permanently remove these outliers from the data
df_model = df.drop(outliers)

```


R2= 0.947325794892571


rmse= 0.09164624386294766

mean of residuals: -7.884104328297002e-16

std of residuals: 0.09167764569489989

19 outliers:

[31, 89, 411, 463, 496, 524, 534, 633, 682, 689, 711, 875, 917, 969, 971, 1299, 1325, 1433, 1454]



```
[86] d_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 1460 entries, 1 to 1460
```

```
Columns: 546 entries, LotFrontage to SaleCondition_Partial
```

```
dtypes: float64(3), int64(50), uint8(493)
```

```
memory usage: 1.3 MB
```

```
[87] d_df.to_csv('./clean_data.csv')
```

```
[ ]
```